



ENTERPRISE AI · DELIVERY PLAYBOOK

Delivering an AI Transformation Project: From Pilot to Production

A practical delivery guide for the program leaders accountable for turning AI ambition into production systems and measurable business results.

Prepared by labwyze

Industry-agnostic edition

Synthesis of analyst, survey, and practitioner sources

Version 1.0

June 2026

Contents

1	Executive Summary	3
2	Why AI Projects Fail Differently	4
3	The Delivery Lifecycle at a Glance	6
4	Phase 1: Readiness & Use-Case Selection	7
5	Phase 2: The Data Foundation Sprint	9
6	Phase 3: Pilot Design & Build	10
7	Phase 4: Production Deployment Is a Separate Project	12
8	Phase 5: Scale & Operate	14
9	The Delivery Team & Budget Shape	15
10	Change Management: Where 70% of the Value Lives	17
11	Governance & Risk in Delivery	18
12	Measuring What Matters	19
13	Pitfalls & Mitigations	20
14	Getting Started: The First 90 Days	21
A	Appendix: Glossary & Sources	22

HOW TO READ THIS PLAYBOOK

This playbook is for the person accountable for delivery: the program leader, transformation lead, or executive sponsor of an AI initiative. It complements the Labwyze AI Center of Excellence playbook, which covers the permanent organizational structure; this volume covers how to deliver a specific transformation project inside or alongside that structure. External claims are numbered to the source list in the Appendix; figures are planning benchmarks, not guarantees.

1 Executive Summary

AI projects fail at roughly twice the rate of conventional IT projects, and the causes are remarkably consistent: inadequate data foundations, pilots designed without a path to production, value measured in technology terms rather than business terms, and change management treated as an afterthought.^{[8][3]} Every one of those causes is a delivery decision, which means every one is preventable.

80%+

Of AI projects fail, roughly twice the rate of non-AI IT projects, per RAND^[8]

95%

Of GenAI pilots fail to deliver measurable ROI, per MIT's State of AI in Business^[4]

70%

Of AI value comes from workforce and process redesign, not the technology, per BCG^[3]

34%

Efficiency gains within 18 months for organizations that follow a structured scaling process, per Gartner^[4]

The activation gap

The 2026 numbers describe a market stuck between ambition and delivery. Deloitte's survey of more than 3,200 executives found 54% of organizations planning to move 40% or more of their AI experiments into production within six months, while only 25% have actually reached that milestone.^[3] McKinsey finds 88% of organizations using AI in at least one function, but only 1% describing their rollouts as mature.^[8] PwC's 2026 CEO survey found 56% of chief executives have realized neither revenue nor cost benefits from AI.^[8] The gap is not a technology problem. It is an execution problem, and it begins in week one.^[3]

The five delivery principles

1. **Sequence strictly; never skip a phase.** Teams that jump from use-case selection to build without auditing data account for the majority of AI projects abandoned in the first six months.^[3]
2. **Treat data readiness as the gate, not a parallel task.** Gartner predicts 60% of agentic AI projects will fail in 2026 specifically because of inadequate data foundations.^{[3][7]}
3. **Design the pilot to answer one question.** Can AI deliver measurable value on this problem, with this data, in this organization? Every design decision serves that question.^[3]
4. **Run production as a separate project.** The companies that scale successfully treat production deployment as its own project with its own resourcing, timeline, and risk plan.^[4]
5. **Hold sponsors accountable for business results, not technology adoption.** The day-90 success criterion is a business metric: cost, cycle time, error rate. Executives measured on technology adoption adopt technology; executives measured on business results build the change structures that deliver them.^[3]

BOTTOM LINE FOR THE SPONSOR

Structured delivery is the difference between the failure statistics and the success ones. Organizations following a structured roadmap reduce failure rates dramatically and report material efficiency and cost gains within eighteen

months.^{[5][4]} The discipline is not exotic: sequence the phases, gate on data, design pilots with a production path, invest the majority of effort in people and process change, and measure business outcomes from day one.

2 Why AI Projects Fail Differently

A leader who has delivered ERP programs, cloud migrations, or large software projects brings valuable discipline to AI delivery, and several instincts that need recalibration. AI projects share the classic failure modes of large IT programs and add several of their own.

2.1 What is the same

Executive sponsorship, scope discipline, realistic timelines, change management, and honest status reporting matter exactly as much as they do on any transformation. The organizations that treat AI as a technology procurement exercise, select a vendor, deploy a model, wait for improvements, get the same result they would get treating an ERP that way: nothing.^[4]

2.2 What is different

Dimension	Traditional IT project	AI transformation project
Requirements	Specifiable up front; the system does what the spec says.	Probabilistic outcomes; performance emerges from data and iteration. ^[5]
The binding constraint	Usually scope, budget, or integration complexity.	Data quality and accessibility. 85% of AI project failures trace to poor data, not poor algorithms. ^[4]
Definition of done	Deployed and stable equals done.	A deployed model is not a finished product; it needs monitoring, retraining, and drift management indefinitely. ^[4]
Where value comes from	Mostly from the system's capabilities.	Roughly 70% from workforce and process redesign; the model is the minority share. ^[3]
Testing	Deterministic; pass or fail against spec.	Statistical; accuracy thresholds, edge-case behavior, drift, and human-review protocols. ^[4]
Completion	A project with an end date.	A permanent capability; the winners are years in and still iterating quarterly. ^[8]

2.3 The two divergence points

Organizations that treat AI as a technology project and those that treat it as a business transformation look identical in week one but diverge sharply by month six.^[3] The two decisions that drive the divergence: where the money goes (the successful ones allocate the majority of investment to people and process change rather than technology^[3]), and what the sponsor is measured on (business results, not adoption metrics^[3]).

FIGURE 2.1 · WHERE AI TRANSFORMATION VALUE ACTUALLY COMES FROM

BCG workforce transformation research: the majority of AI value comes from how work is redesigned around the capability, not from the capability itself.^[3]


THE UNCOMFORTABLE STATISTIC TO PLAN AGAINST

More than 80% of AI projects fail, twice the rate of non-AI IT projects.^[8] This is not a reason to avoid AI; it is a reason to plan differently. The failure causes are known, consistent, and structural, which means a delivery approach designed against them changes the odds materially. That is what the rest of this playbook does.

3 The Delivery Lifecycle at a Glance

A reliable AI delivery lifecycle has five phases, each producing a specific deliverable that the next phase depends on. The single most important rule: no phase begins before its predecessor is complete.^[3]

FIGURE 3.1 · THE FIVE-PHASE AI DELIVERY LIFECYCLE



3.1 Timelines and what they depend on

Scope	Indicative timeline	What stretches it
First pilot to result	3 to 4 months for a well-selected use case; 90 days for the first structured sequence. ^{[1][3]}	Data readiness discovered late; unclear success metrics; scope creep.
Focused initiative	6 to 12 months to production value. ^[5]	Integration complexity; change-management underinvestment.
Enterprise transformation	12 to 24 months to scaled maturity. ^{[5][6]}	Governance gaps; skills shortages; portfolio sprawl without prioritization.

THE PHASE-SKIPPING TRAP

Teams that skip ahead, especially those that jump from use-case selection to build before auditing data, account for the majority of AI projects abandoned in the first six months.^[3] The lifecycle is not bureaucracy; each phase exists because its absence is a documented, recurring cause of failure. The discipline of sequencing is the cheapest risk control available.

4 Phase 1: Readiness & Use-Case Selection

The first phase determines the odds of everything after it. Its two products: an honest readiness picture across data, infrastructure, skills, and business alignment, and a use case selected on evidence.

4.1 The readiness assessment

Evaluate the organization's current state across four dimensions: data maturity, technical infrastructure, team capabilities, and business alignment. The purpose is to identify the gaps that would derail implementation later, while they are still cheap to fix.^[1] The deliverable is a readiness report with a gap analysis; typical duration is two to four weeks for smaller organizations and four to six for enterprises.^[1] A useful discipline is to score each dimension on a one-to-five maturity scale and require a minimum threshold before proceeding to pilots.^[1]

4.2 Selecting the use case

The selection criteria that predict success are consistent across sources: data availability, business impact potential, technical complexity, and stakeholder support.^[1] A good first use case addresses a specific pain point with measurable outcomes achievable within three to four months.^[1]

Criterion	What to verify before committing
Business impact	A named business owner; a pain point with a dollar, hour, or error-rate figure attached; a metric leadership already tracks.
Data availability	The data exists, is accessible, and is of workable quality, verified in Phase 2, but screened here.
Technical complexity	Achievable with proven patterns; not dependent on frontier capabilities or unproven integrations.
Stakeholder support	The affected team wants the solution; a skeptical process owner sinks more pilots than a weak model.
Path to production	A defined scaling route exists before the pilot starts; pilots without one are experiments.

BALANCE THE PORTFOLIO, SEQUENCE THE DELIVERY

At program level, a balanced portfolio mixes foundational, differentiated, and strategic initiatives.^[2] But delivery is sequential: prove the model on one well-chosen use case before running many in parallel. The strongest strategic filter is differentiation: AI that amplifies what the organization uniquely knows creates advantage, while AI identical to every competitor's tool bundle institutionalizes parity.^[8]

5 Phase 2: The Data Foundation Sprint

The most common reason AI pilots fail is not poor AI performance; it is inadequate data.^[3] Gartner finds 63% of organizations lack the data management practices required for AI, and predicts 60% of agentic AI projects will fail in 2026 specifically because of inadequate data foundations.^{[3][7]} This phase is the gate that keeps your project out of those statistics.

5.1 A targeted sprint, not an enterprise program

The data audit at this stage is deliberately narrow: a targeted two-week sprint answering four questions about the specific use case selected in Phase 1.^[3]

1. **Where does the relevant data live?** Map what is in the ERP, what is in operational systems, what is in spreadsheets and shared drives, and what does not exist yet.^[3]
2. **How much of it exists?** Volume and history sufficient for the intended approach.
3. **How clean and consistent is it?** Completeness, duplicates, referential integrity, and consistency across sources.
4. **How quickly can the technical team access it?** Access paths, permissions, latency, and security constraints for development.^[3]

5.2 The gate decision

The sprint ends in one of three outcomes: proceed (data is workable), remediate first (a scoped data-fix effort with its own timeline before build begins), or reselect (the use case is not viable with available data; return to Phase 1 with the evidence). All three are successes; the failure is proceeding on hope. Data maturity below a minimum viable threshold produces models that deliver inaccurate insights or fail outright.^[7]

WHY THIS PHASE PAYS FOR ITSELF

Data discoveries made in a two-week sprint cost two weeks. The same discoveries made mid-build cost the build. Evaluate data maturity across availability, lineage, integration, accuracy, latency, security, and governance for the target use case, and let the evidence set the plan.^[7]

6

Phase 3: Pilot Design & Build

A well-designed AI pilot answers one specific question: can AI deliver measurable value on this problem, with this data, in this organization? Every design decision should serve that question, and nothing else.^[3]

6.1 Design before build

- **Success metrics first.** Define the business metric, its baseline, and the threshold that constitutes success before the first sprint. A pilot without a pre-agreed success definition ends in a dispute, not a decision.^{[1][3]}
- **Fallback strategy.** Agree what happens if the pilot misses: iterate, pivot, or stop. Failure criteria are as important as success criteria.^[1]
- **Scope guardrails.** A lightweight governance structure preventing the three most common pilot failure modes: scope creep, accountability gaps, and result disputes.^[3]
- **The business owner role.** A named business owner accountable for the outcome, with the capability development that role requires; it demands judgment most traditional managers have not needed before.^[3]

6.2 The build rhythm

Development follows agile methodology with two-week sprints. Key milestones: data pipeline establishment, model development and validation, system integration, and user acceptance testing. Plan for two to three iteration cycles before reaching target performance; a model that hits its thresholds on the first pass is the exception, not the plan.^[1] A cross-functional team of four to six covers the typical first pilot.^[1]

6.3 Prototype early, decide cheaply

The strongest 2026 delivery pattern replaces long advisory cycles with working prototypes built early, pressure-testing architecture, data pipelines, and governance controls before large budgets are committed.^[7] The prototype is not the product; it is the cheapest possible way to surface the problems that would otherwise surface in production.

MEASURE AT THE BUSINESS-UNIT LEVEL, NOT THE PILOT LEVEL

One of the most insidious pitfalls: initiatives that look successful in a controlled environment while delivering no measurable P&L impact at all.^[4] Design the measurement so the pilot's result is read in business-unit numbers, cost, cycle time, error rate, revenue, not in a sandbox metric that flatters the model.

7

Phase 4: Production Deployment Is a Separate Project

The companies that scale most successfully treat production deployment as a separate project from the pilot, with its own resourcing, timeline, and risk management plan.^[4] A pilot that works is evidence, not a product. The distance between them is where most successful pilots quietly die.

FIGURE 7.1 · WHAT CHANGES BETWEEN PILOT AND PRODUCTION

The pilot proved	Production requires
• The model performs on the target problem • The data supports the approach • Users can work with the output • The business metric moves • The organization can deliver AI at small scale	• Hardened integration with core systems and security • Monitoring infrastructure that flags drift before it becomes a quality problem^[4] • Escalation protocols routing out-of-distribution cases to human review automatically^[4] • Support model, incident response, and a named operational owner • Retraining cadence and model lifecycle management^[4]

7.1 The production checklist

Requirement	What good looks like
Monitoring & drift detection	Automated performance monitoring with thresholds; drift flagged before users notice quality decay. ^[4]
Human escalation	Out-of-distribution and low-confidence cases routed to a human reviewer automatically rather than letting the model guess. ^[4]
Operational ownership	A named owner for the system in production; accountability does not dissolve at go-live.
Retraining & lifecycle	A defined cadence and trigger conditions for retraining; a deployed model is not a finished product. ^[4]
Security & compliance	Access controls, audit logging, and regulatory alignment carried from governance design into the production stack.
Feedback loops	Real-time performance data feeding continuous refinement of the model and the process around it. ^[6]

THE QUIET FAILURE MODE

Ongoing oversight, retraining, and governance is the phase most roadmaps treat as an afterthought, and it is where many otherwise successful programs quietly lose momentum.^[4] The gains reported by structured scalers, 34% efficiency and 27% cost reduction within eighteen months, do not materialize if the monitoring and feedback loops are not in place.^[4]

8

Phase 5: Scale & Operate

Scaling is not running the pilot bigger. It is converting one proven pattern into organizational capability: replicating what worked, industrializing the plumbing, and running AI as a permanent operation rather than a project.

8.1 The scaling moves

- **Replicate patterns, not projects.** Extract what made the pilot work, the data pipeline design, the governance protocol, the escalation model, the measurement approach, into reusable assets the next use case starts from. Patterns that prove reliable get reused rather than rebuilt.
- **Interoperable ecosystems.** Ensure AI tools work across the technology stack rather than creating islands of automation that do not communicate, which breed technical debt and security exposure.^[6]
- **Portfolio discipline.** Prioritize the next use cases with the same evidence-based selection as the first; resist chasing every available use case.^[9]
- **Operating model.** Define who owns models, who manages pipelines, how teams collaborate, and which roles must evolve or be added. AI at scale is an operational and cultural transformation, not only a technical one.^[7]

8.2 Transformation has no completion date

A pilot can be marked done; a transformation cannot, because the technology and competitive dynamics keep shifting.^[9] The organizations winning at AI are years into their transformations and still iterating every quarter.^[8] Deliver the project with an end date; build the capability without one. Practically, this means the delivery project's final deliverable is the operating rhythm that continues after it: monitoring, retraining, quarterly portfolio review, and continuous feedback loops.^[6]

9 The Delivery Team & Budget Shape


Two structural choices set the delivery project's cost realism: who is on the team, and where the money goes. Both have well-documented shapes that surprise leaders coming from traditional IT delivery.

9.1 The cross-functional core

AI delivery requires machine learning engineers, data engineers, product leaders, domain subject-matter experts, and governance expertise aligned under a shared operating model.^[7] The first-pilot team is typically four to six people cross-functionally^[1]; what matters more than headcount is that business stakeholders work alongside the engineers throughout, rather than receiving a handoff at the end.^[7] Skills are a real constraint: PwC found skill gaps among the top three barriers to scaling AI, ranking above funding and tooling.^[7] Most mid-market organizations should not try to build a full internal AI capability from scratch; the experienced-talent market is extremely competitive and building internally takes years the business may not have. Strategic partnerships to fill engineering and data science gaps are a legitimate acceleration move, not an admission of weakness.^{[4][6]}

9.2 Where the budget actually goes

FIGURE 9.1 · TYPICAL AI TRANSFORMATION BUDGET DISTRIBUTION



Indicative distribution synthesized from implementation guides; pilots commonly run \$50K to \$200K, enterprise transformations \$500K to \$5M+. ^[5]

THE BUDGET LINE TO CHALLENGE

The conventional 10% change-management line sits in tension with the finding that roughly 70% of AI value comes from workforce and process redesign.^[3] Deloitte's research makes the same point from the other side: AI spending keeps rising, but only a small fraction goes toward the people and process changes that determine whether the technology gets used.^[4] Treat the standard distribution as a floor for change investment, not a ceiling; in hands-on operating environments the human adoption curve, not the model, is the rate-limiting factor.^[4]

10

Change Management: Where 70% of the Value Lives

If 70% of AI value comes from workforce and process redesign,^[3] then change management is not a support function of the delivery project; it is the delivery project. The model is the enabler. The redesigned way of working is the product.

10.1 What AI change management involves

Workstream	What it means in an AI delivery
Process redesign	Redesigning the workflow around the AI capability, not bolting AI onto the old process. This is where the majority of the value is created or lost. ^[3]
Role evolution	Defining how affected roles change: what the human now decides, what the system handles, and where judgment enters. New oversight roles emerge. ^[7]
Training & AI literacy	Employees need training to work alongside AI effectively, timed close to go-live and role-specific rather than generic. ^[6]
Trust building	Escalation protocols and human-review paths are adoption tools as much as risk controls; users trust systems they can override and understand.
Communication	Sponsor-led, honest about what changes for whom, and sustained past go-live into the operating rhythm.

10.2 The adoption curve is the rate limiter

In manufacturing, logistics, and other hands-on operating environments, the human adoption curve, not the model, is the rate-limiting factor on value.^[4] The practical consequence for delivery planning: schedule value realization against adoption milestones, not deployment milestones. A deployed system with 20% adoption delivers roughly 20% of its case.

THE SPONSOR'S ACCOUNTABILITY TEST

The cleanest structural fix is measurement: hold the executive sponsor accountable for business results at day 90, cost reduction, cycle time, error rate, not for technology adoption.^[3] Sponsors measured on business results build the change structures required to deliver them; sponsors measured on adoption ship tools.^[3]

11 Governance & Risk in Delivery

For a first delivery, the goal is not a mature enterprise governance framework; it is a functional, lightweight structure that prevents the three most common failure modes: scope creep, accountability gaps, and result disputes.^[3] Only 20% of enterprises have mature AI governance,^[3] and waiting for maturity before delivering is its own failure mode.

11.1 The minimum viable governance set

- **Decision rights.** Who approves scope changes, who calls the pilot result, and who owns the go/no-go to production. Written down before the build starts.
- **Risk classification.** Assess the use case's risk tier early, especially for anything touching individuals' finances, employment, or regulated decisions, where EU AI Act high-risk obligations may apply.
- **Model risk controls.** Fairness, transparency, data privacy, security, and model risk managed proportionate to the use case, and embedded from the roadmap rather than retrofitted, so systems are compliant by design.^[7]
- **Audit trail.** Logging of model decisions and human overrides from the pilot onward; cheaper to build in than to bolt on.
- **Escalation & human oversight.** The human-review protocol designed in Phase 4 doubles as the governance control regulators and auditors ask about first.

WHERE SOURCES DIFFER: AGENTIC AMBITION VERSUS GOVERNANCE REALITY

Vendor enthusiasm for autonomous agents runs ahead of governance readiness: Gartner predicts over 40% of agentic AI projects will be canceled by the end of 2027, and 60% will fail in 2026 on data foundations alone.^{[8][7]} The pragmatic posture: build governance frameworks that can accommodate autonomous agents, but do not restructure the organization around capabilities that are not production-ready yet.^[8] Invest in governance infrastructure that can absorb new regulatory requirements without starting over each time a rule lands.^[8]

12 Measuring What Matters

Measurement discipline separates the organizations that scale from those that stall: mature adopters regularly conduct financial and ROI analysis on their AI projects, building the case for continued investment as they go.^[2] The design principle is to measure outcomes rather than activity.^[9]

12.1 A three-layer metric stack

Layer	Representative metrics	Who reads it
Business outcomes	Cost reduction, cycle-time improvement, error-rate reduction, revenue impact, measured at business-unit level against a pre-pilot baseline. ^{[3][4]}	Sponsor and steering; the day-90 success criterion. ^[3]
Adoption & process	Share of in-scope volume through the new process; active usage; override and escalation rates; training completion.	Business owner and change lead; the leading indicator of the outcome layer.
Technical health	Accuracy against thresholds, drift signals, latency, incident counts, cost per inference.	Delivery and operations teams; the early-warning layer.

12.2 Measurement rules that prevent disputes

- **Baseline before build.** Capture the pre-pilot baseline in Phase 3 design; a result without a baseline is an anecdote.
- **Business-unit level, not sandbox level.** Pilots can look successful in controlled environments while delivering no P&L impact; measure where the money is.^[4]
- **Business metrics for sponsors, technical metrics for teams.** The success criterion at day 90 is a business number, not model accuracy.^[3]
- **Review, do not just report.** Scale what works, adjust or retire what does not, and feed the evidence into the next selection round.^[9]

13 Pitfalls & Mitigations

The AI delivery failure modes are consistent enough across the research to read as a checklist. Each has a structural mitigation already embedded in the lifecycle.

Pitfall	How it shows up	Mitigation
Technology procurement framing	Select vendor, deploy model, wait for improvement; nothing happens. ^[4]	Frame as business transformation; majority of investment to people and process change. ^[3]
Phase skipping	Use case chosen and build started before the data audit; project abandoned within six months. ^[3]	Hard gate: no build before the data foundation sprint completes. ^[3]
Inadequate data	Model underperforms or fails; 85% of failures trace to data quality. ^[4]	Two-week targeted data sprint with a proceed / remediate / reselect decision. ^[3]
Pilot purgatory	Pilots demonstrate possibility but never reach production; executive confidence erodes. ^[2]	Path to production defined before the pilot starts; production run as its own project. ^[4]
Sandbox success	Pilot metrics look great; no P&L impact appears. ^[4]	Measure at business-unit level against a pre-pilot baseline. ^[4]
Change as afterthought	Tool ships; adoption stalls; value never materializes. ^[4]	Change investment as the majority workstream; value scheduled against adoption milestones. ^{[3][4]}
Sponsor measured on adoption	Technology gets adopted; business results do not appear. ^[3]	Day-90 success criterion is a business metric owned by the sponsor. ^[3]
Islands of automation	Disconnected tools that do not communicate; technical debt and security exposure. ^[6]	Interoperability standards; portfolio-level architecture review. ^[6]
Post-go-live neglect	No monitoring or retraining; the program quietly loses momentum. ^[4]	Monitoring, drift detection, escalation, and retraining cadence as production entry criteria. ^[4]

14 Getting Started: The First 90 Days

A first AI transformation initiative can be sequenced across ninety days, with each phase producing the deliverable the next depends on.^[3] The point of the ninety days is not to finish the transformation; it is to produce one measured business result and a repeatable delivery pattern.

FIGURE 14.1 · THE 90-DAY DELIVERY SEQUENCE



14.1 The first five moves

1. **Name the sponsor and the business metric.** One accountable executive, one day-90 business number.^[3]
2. **Run the readiness assessment.** Four dimensions, honest scoring, gaps identified while they are cheap.^[1]
3. **Select one use case on evidence.** Impact, data, complexity, stakeholder support, and a defined path to production.^[1]
4. **Gate on the data sprint.** Two weeks, four questions, three honest outcomes.^[3]
5. **Fund the change, not just the build.** Process redesign and adoption as the majority workstream from day one.^[3]

HOW LABWYZE CAN HELP

labwyze supports executives and program leaders in delivering AI transformation: readiness assessment and evidence-based use-case selection, data foundation review, pilot design with defensible success metrics, production and operating-model planning, and the governance and change structures that turn a working model into a business result. The goal is a first delivery that produces a measured outcome in ninety days and a repeatable pattern the organization owns afterward.

This playbook is a synthesis of public sources current to mid-2026 and is intended as planning guidance, not legal or investment advice. Figures are benchmarks that vary by organization size, industry, and maturity.

A

Appendix: Glossary & Sources

A.1 Glossary

Activation gap

The distance between AI experiments and production deployments; the defining enterprise AI problem of 2026.

Data foundation sprint

A targeted two-week audit of the selected use case's data across location, volume, quality, and access, ending in a proceed / remediate / reselect gate.

Drift

Degradation of model performance as real-world data shifts away from training data; the reason a deployed model is never a finished product.

Escalation protocol

The automatic routing of out-of-distribution or low-confidence cases to human review rather than letting the model guess.

Islands of automation

Disconnected AI deployments that do not interoperate, creating technical debt and security exposure.

Pilot purgatory

The stalled state where pilots prove technical possibility but never reach production value.

Sandbox success

A pilot that meets its metrics in a controlled environment while producing no measurable P&L impact.

Out-of-distribution

Inputs materially different from what the model was trained on, where its outputs are least reliable.

MLOps

The engineering discipline of deploying, monitoring, retraining, and managing machine-learning systems in production.

Day-90 criterion

The pre-agreed business metric (cost, cycle time, error rate) by which the first delivery is judged, owned by the executive sponsor.

A.2 Sources

Sources numbered as cited inline. Survey and analyst findings (Deloitte, Gartner, McKinsey, PwC, BCG, RAND, MIT, IBM) are reported via the listed publications; figures are planning benchmarks synthesized across sources.

1. Space-O, "AI Implementation Roadmap: 6-Phase Guide for 2026" (2026). Readiness dimensions; pilot selection criteria; sprint rhythm; team size; timelines. spaceo.ai
2. AI Assembly Lines, "AI Transformation Roadmap 2026" (2026). Four value phases; portfolio balance; pilot purgatory; ROI analysis discipline. aiassemblylines.com
3. AI Assembly Lines, "How to Start an AI Transformation in 2026: A 5-Phase Framework" (Apr 2026). Deloitte activation gap; Gartner data findings; BCG 70% value; day-90 business criterion; phase discipline; lightweight governance. aiassemblylines.com
4. AI Assembly Lines / Assembly, "AI Transformation Roadmap: The 6-Phase Guide" (2026). MIT 95% pilot failure; Gartner 85% data-quality failures; production as separate project; monitoring and escalation; 34%/27% structured-scaling gains; change-management underinvestment. aiassemblylines.com

5. GrowExx, "AI Implementation Roadmap: Complete 2026 Strategy Guide" (Apr 2026). Budget distribution; cost ranges; 70-85% failure reduction with structured roadmaps; timelines by size. growexx.com
6. The Ksquare Group, "AI Transformation Roadmap: 5 Powerful Steps" (Apr 2026). Islands of automation; siloed thinking; feedback loops; partnerships; 12-18 month maturity. theksquaregroup.com
7. RTS Labs, "Enterprise AI Roadmap: The Complete 2026 Guide" (May 2026). Gartner 60% agentic failure on data; PwC 34% financial impact and skills-gap findings; prototype-early methodology; operating model; compliance by design. rtslabs.com
8. Larridin, "The Complete Guide to AI Transformation (2026)" (May 2026). RAND 80%+ failure; BCG 5% at scale; McKinsey 1% maturity; PwC CEO survey; Gartner 40% agentic cancellations; differentiation principle; permanent capability. larridin.com
9. USAII, "The AI Transformation Roadmap Every Leader Needs in 2026" (2026). IBM CAIO trend; KPMG disconnected-projects finding; outcomes-over-activity; annual strategy review. usaii.org

A NOTE ON EVIDENCE QUALITY

Failure-rate statistics vary by definition and source (RAND's 80%+, MIT's 95% for GenAI pilots, and Gartner's domain-specific predictions measure different things); they agree on direction and rough magnitude, not on a single number. Budget distributions, timelines, and gain figures are indicative benchmarks from implementation guides, several published by consultancies with services to sell. Use them as planning anchors and validate against your own baseline; the playbook's structural recommendations rest on the points where independent sources converge.