



FINANCE TRANSFORMATION · AI READINESS
PLAYBOOK

Agentic AI in the Finance Function: A Readiness Playbook

A decision-grade guide for CFOs and finance leaders moving from AI experiments to governed, value-generating autonomous finance in 2026.

Prepared by labwyze

Industry-agnostic edition

Synthesis of SAP, Big 4, analyst, and regulatory sources

Version 1.0

June 2026

Contents

1	Executive Summary	3
2	Why Now: The 2026 Inflection Point	4
3	What "Agentic" Actually Means for Finance	6
4	The Finance AI Maturity Model	8
5	Where Agents Create Value: Use Cases by Function	10
6	The Data Foundation: Why Readiness Is the Real Gate	13
7	Governance, Risk & the EU AI Act	15
8	The Business Case: ROI, Payback & KPIs	18
9	Buy, Build, or Embed: The Sourcing Decision	20
10	A 90-Day Readiness Roadmap	22
11	Risks, Pitfalls & Mitigations	24
12	Getting Started	25
A	Appendix: Glossary & Sources	26

HOW TO READ THIS PLAYBOOK

Statements drawn from external research are numbered inline to the source list in the Appendix (for example, [5]). Where vendor framing and independent or regulatory views diverge, that divergence is flagged in a teal "Where sources differ" box. Synthesis and recommendations that are the author's own appear without a source marker. Figures are planning benchmarks, not guarantees for any specific organization.

1 Executive Summary

In 2026 the question facing finance leaders is no longer whether to adopt AI, but how to move from scattered experiments to governed, autonomous operations that produce measurable returns. The technology has crossed from copilots that suggest to agents that act, and boards are now asking for results rather than pilots.^{[1][6][13]}

54%

Of CFOs name integrating AI agents a digital-transformation priority for 2026^{[2][7]}

4–9 mo

Typical payback for a focused task-automation agent (for example, accounts payable)^[3]

86%

Of finance functions saw no significant ROI from earlier AI in 2024, for lack of strategy and targeting^[8]

Aug 2026

EU AI Act high-risk obligations become enforceable^{[17][19]}

What has changed

Generative AI tools require a human to trigger them and produce output for a human to review. Agentic systems perceive context, plan multi-step actions, and execute end-to-end within defined rules and approval thresholds.^[3]
^[13] In finance this is a shift on the order of moving from spreadsheets to ERP: the numbers no longer just update, they trigger action.^[9] Early adopters report close times cut by up to half, collections transformed, and forecasting that refreshes continuously rather than monthly.^{[8][3]}

The decisions leadership must own

1. **Ambition and maturity.** Decide honestly where the function sits today on the maturity curve, and how far it intends to move in twelve months (Section 4).
2. **Sequence of use cases.** Choose where to start based on data readiness, volume, and the cost of errors, not on vendor enthusiasm (Section 5).
3. **Governance posture.** Decide the guardrails, oversight model, and EU AI Act response before the first agent goes live, not after (Section 7).
4. **Sourcing.** Buy to accelerate standard processes, build to differentiate, and decide deliberately which is which (Section 9).

Key takeaways

- **Data readiness is the binding constraint, not model quality.** The agents are production-ready; the data and the governance usually are not. Organizations that skip the data and governance groundwork hit expensive rebuilds.^{[4][8]}
- **Start where the return is documented.** Accounts payable is the proving ground for most finance functions: structured data, high volume, clear controls, and the fastest measured payback.^{[13][3]}

- **Governance is an accelerator, not a brake.** CFOs who build governance in from the start see faster adoption and stronger regulatory confidence than those who treat it as an afterthought.^{[3][13]}
- **The regulatory clock is real.** For organizations with EU exposure, high-risk AI obligations are enforceable from August 2026, and a credible readiness program takes the better part of a year.^{[17][19]}

BOTTOM LINE FOR THE BOARD

If the finance function has not completed an AI readiness assessment, mapped its data foundation, and stood up a governance framework, that is the action to authorize this quarter. The competitive risk of not deploying now exceeds the operational risk of a well-governed start, and the gap between early movers and evaluators is already widening.^{[6][3]}

2 Why Now: The 2026 Inflection Point

Three forces have converged to make 2026 the year finance AI moves from optional to expected: the technology became production-ready, the economics became documented, and the board became impatient.

2.1 From experiments to expectations

For two years finance teams ran proofs of concept. Many produced insight but little return: industry analysis found roughly 86% of finance functions reported no significant ROI from AI in 2024, largely for want of a clear strategy and targeted use cases.^[8] That patience has run out. Boards and chief executives now expect production results, and nearly half of CFOs report direct pressure from leadership to deploy AI across operations.^[13] The defining 2026 sentence from vendors and operators alike is blunt: AI for its own sake is a waste; the point is outcomes embedded in the workflow.^[13]

2.2 The autonomy shift

The capability that changed the calculus is autonomy. Earlier automation followed scripts: rule-based robotic process automation, deterministic macros, and chatbots. Agentic systems blend reasoning models, policy libraries, and system access so software can observe, decide, and act in real time across finance systems.^[9] The practical difference for a CFO is that an agentic system does not merely surface an anomaly for a human to investigate; it investigates, resolves, escalates if needed, and documents the full process without manual coordination.^[3]

THE RISK CALCULUS HAS INVERTED

By 2026 the competitive risk of not deploying exceeds the operational risk of deploying for many finance leaders: a talent-acquisition risk as professionals prefer technology-forward employers, and a scalability risk as transaction volumes grow faster than headcount budgets.^[6] Agentic AI lets a finance function handle several times the transaction volume without several times the headcount, while the team it has focuses on analysis and judgment.^[5]

2.3 The trajectory is steep

Independent forecasts point the same direction. Gartner projects that at least 15% of day-to-day finance work decisions will be made autonomously by AI by 2028, and that the large majority of finance analytics will be automated within a few years.^{[5][9]} A Deloitte-cited survey found roughly 80% of finance and accounting professionals expect AI agents to be standard tools within five years.^[5] The window in which early adopters gain a structural advantage over those still evaluating is already closing.^[5]

2.4 The ERP vendors are moving in

The shift is not only a market of point solutions; it is being built into the core ERP. SAP has framed an Autonomous Enterprise vision in which assistants and agents work alongside people, with Autonomous Finance

combining assistants and agents to give CFOs more insight and control. A Financial Closing Assistant, for example, is designed to orchestrate the full close lifecycle across posting, accruals, journal validation, error resolution, and intercompany reconciliation, with general availability of several agentic capabilities planned through the back half of 2026.^[1] For organizations on or moving to S/4HANA, this means the AI conversation and the platform conversation are now the same conversation.

WHERE SOURCES DIFFER: HYPE VERSUS REALITY

Vendors describe production-ready autonomy; independent commentators urge a reality check. The durable framing is that agentic AI is an amplifier of finance leadership, not a replacement for it: it frees analysts from preparation work, but strategic oversight remains human.^[15] A CFO's instinct for incremental pilots with clear failure criteria is, in this case, an asset rather than a brake, because autonomy cuts both ways, faster decisions but also faster mistakes if the guardrails are wrong.^[15]


3 What "Agentic" Actually Means for Finance

Precision of language matters here, because the word "agent" is used loosely and the difference determines your governance burden, your data needs, and your risk. The useful distinction is between predefined workflows and genuinely autonomous agents.^[12]

3.1 The autonomy spectrum

Finance capability sits on a spectrum from scripted automation to self-optimizing systems. Most finance functions today are between the first and second stages and are deciding how far right to move, and how fast.

FIGURE 3.1 · THE FINANCE AUTOMATION AUTONOMY SPECTRUM



More autonomy on the right means more value and more governance burden. The right target is process-specific, not a single setting for the whole function.

3.2 What makes an agent different from a bot

What separates agentic AI from older automation is that agents act and decide within context. They read governance rules, historical resolutions, and live signals, and act only when a move aligns with policy and probability thresholds.^[9] Where a script breaks on an exception, an agent reasons through it or escalates it. Where robotic process automation needs a human to define every path, an agent plans the path. This is why agentic systems are described as the first automation that can handle the full complexity of real finance workflows, including the exceptions and the cross-system dependencies, without a human managing every deviation.^[5]

3.3 The four layers under every finance agent

Finance technologists will recognize the enabling architecture beneath a credible agent.^[9]

Layer	What it does	Why it matters to finance
Unified data fabric	Transaction-level data from ERP, planning, and treasury feeds the agent a single source of truth.	This is the control tower; without it, agents reason on conflicting numbers.
Reasoning layer		

Layer	What it does	Why it matters to finance
	Finance-tuned models interpret accounting guidance, internal policy, and history to choose the next best action.	Determines whether the agent follows your controls or improvises around them.
Action layer	Secure connections that let the agent read and write to core systems through APIs.	Turns insight into posted entries, payment runs, and reconciliations.
Governance layer	Authorization limits, escalation criteria, immutable logs, and drift monitoring.	The difference between a controlled agent and an uncontrolled liability.


A DELIBERATELY CONSERVATIVE DEFAULT

Because finance is a tightly regulated, high-stakes domain, workflow-style automations that offer control and predictability tend to be adopted in practice before highly autonomous, low-predictability agents.^[12] The pragmatic path is to start with bounded agents under human approval and earn the right to greater autonomy with each proven, audited cycle.

4 The Finance AI Maturity Model

A maturity model turns an abstract ambition into an honest baseline and a sequence of moves. The Labwyze model places a finance function on one of five levels and makes explicit what changes at each step: not just the technology, but the data, the operating model, and the governance.

FIGURE 4.1 · THE LABWYZE FINANCE AI MATURITY MODEL



4.1 What actually changes at each level

Level	Technology	Data & operating model	Governance
1 Manual	Spreadsheets, manual reconciliation.	Fragmented data; no common chart of accounts discipline.	Manual controls only.
2 Assisted	RPA bots, GenAI copilots for drafting and search.	Some standardization; humans still own every step.	Ad hoc; AI use largely ungoverned.
3 Orchestrated	Bounded agents for specific workflows under approval.	Clean data for target processes; redesigned hand-offs.	Authorization limits and logging for live use cases.
4 Autonomous	Agents execute end-to-end; exception-based oversight.	Unified data fabric across core finance; new roles (agent supervisors).	Full lifecycle risk management and drift monitoring.
5 Self-optimizing	Multi-agent orchestration; continuous processes.		

Level	Technology	Data & operating model	Governance
		Real-time data across functions; finance reorganized around oversight.	Continuous controls and assurance, audit-ready by design.

HOW TO USE THE MODEL

Do not target Level 5 everywhere. Target a realistic move of one to two levels over twelve months, process by process. A function can be at Level 4 in accounts payable and Level 2 in financial planning at the same time, and that is the correct, sequenced state. The model is a planning instrument, not a scoreboard.

4.2 The honest caveat on maturity

Maturity is not bought; it is earned through data work and governance. The most common failure is an organization that purchases a high-autonomy capability while its data and controls sit at Level 2, then blames the technology when results do not appear. The level you can operate safely is set by your data foundation and your governance, not by the sophistication of the agent you license.

5

Where Agents Create Value: Use Cases by Function

Value is uneven across the finance function. The right first moves combine high transaction volume, structured data, a clear cost of error, and a control environment that an agent can operate inside. The matrix below scores the major domains on the dimensions that decide sequencing.

FIGURE 5.1 · USE-CASE VALUE AND READINESS MATRIX

Finance domain	Value potential	Typical data readiness	Autonomy fit today	Suggested sequence
Accounts payable / invoice-to-pay	High	High	Strong	Start here
Accounts receivable / collections	High	Medium	Strong	Wave 1–2
Record-to-report / close	High	Medium	Growing	Wave 2
Controls, audit & compliance	High	Medium	Growing	Cross-cutting
Treasury & liquidity	Med–High	Medium	Growing	Wave 2–3
FP&A / forecasting & scenarios	High	Variable	Augment	Wave 2–3
Tax & statutory reporting	Medium	Variable	Selective	Later

5.1 Accounts payable: the proving ground

Accounts payable is where most finance functions should begin. Around 72% of finance leaders see it as the obvious starting point, because it combines structured data, scale, control, and accountability in a way few finance processes do.^[13] Agents verify invoices against contracts and purchase orders, detect duplicates and anomalies, route exceptions for approval, and execute scheduled payment runs.^[19] Reported outcomes include large reductions in invoice-processing cost and cycle time, and autonomous agents have delivered an average return well above general AI projects in this domain.^{[11][13]}

5.2 Accounts receivable: a direct line to cash

Collections and cash application have a revenue impact that is measurable within months. Agents match incoming payments to open invoices from remittance data, prioritize high-risk accounts, prompt customers, and resolve simple disputes, reducing days sales outstanding and unapplied cash.^{[10][3]}

5.3 The close: toward a continuous process

In record-to-report, agents orchestrate posting, accruals, journal validation, reconciliation, and intercompany matching, surfacing bottlenecks and resolving discrepancies in real time.^[1] The destination is a continuous close in which assurance is a minute-by-minute process rather than a month-end scramble, and early adopters report close times cut by up to half.^{[8][9]}

5.4 FP&A and treasury: augment first, automate carefully

In planning, agents refresh forecasting models in real time, pull data from sales, supply chain, and macro sources, and run scenarios on demand rather than on a quarterly schedule, with teams reporting a flip in the ratio of time spent on analysis versus preparation.^{[9][3]} In treasury, agents monitor intraday liquidity, trigger sweeps or short-term placements within policy, and flag cash-flow anomalies that monthly reconciliation would miss.^[9] A common and sensible sequencing pattern is to prove value in payables and receivables, then extend the same governance to treasury.^[3]

THE SEQUENCING PRINCIPLE

Sequence by where manual effort is highest, error rates are most costly, and clean data already exists.^[3] Accounts payable usually offers the fastest, most documented return and is the recommended first move; receivables follow for direct cash impact; the close and treasury extend the same guardrails; planning is most often an augmentation rather than full autonomy.

6

The Data Foundation: Why Readiness Is the Real Gate

The single most reliable predictor of whether a finance AI program succeeds is the state of its data. Agentic systems are only as effective as the data they receive, and the organizations that bypass data readiness consistently encounter expensive rebuilds and delayed timelines.^{[8][13]}

6.1 What "ready" data means for an agent

- **A standardized chart of accounts.** Inconsistent account structures force agents to guess; standardization is the foundation.^[8]
- **Unified entity and vendor mappings.** Duplicate or mismatched master data is the classic and expensive failure mode, and agents amplify it at machine speed if it is not fixed first.^[8]
- **Consistent tagging and reference data.** Reliable categorization lets an agent distinguish a legitimate anomaly from an error without a human in the loop.^[13]
- **Accessible, transaction-level data.** A unified data fabric across ERP, planning, and treasury is what gives agents a single source of truth.^[9]

6.2 The structured-data advantage

This is also why accounts payable works so well as a first use case: it runs on structured data, where high-quality history lets a system tell legitimate exceptions from mistakes without constant oversight.^[13] The corollary is a warning. Processes that depend on unstructured or fragmented data are harder to automate safely, which is one reason planning and tax sit later in most sequences.

WHERE SOURCES DIFFER: FIX-FIRST VERSUS LEARN-AS-YOU-GO

Some vendors suggest agents can be deployed quickly and will improve as they learn. Independent practitioners are firmer: the foundation must come first, because agents trained or operated on poor data institutionalize the errors. The defensible position for a CFO is to treat a focused data-readiness baseline as a prerequisite for the first production agent, not a parallel workstream to be caught up later.^{[8][13]}

6.3 The migration connection

For organizations also moving to S/4HANA, the data question now reframes the migration itself. The leading question in 2026 is no longer only what AI can do after go-live, but what to fix, harmonize, or redesign during the migration so that AI works afterward.^[24] Data transformation, system redesign, and integration become building blocks of AI adoption rather than technical steps on the way to a new ERP. Clean-core discipline, keeping the digital core standard and pushing extensions to a platform layer, is part of the same story: a lean, standard core is what lets agentic innovation be adopted quickly.^[24]

6.4 A practical readiness baseline

Before any agent is licensed, a finance function should be able to answer four questions with evidence: Is the chart of accounts standardized and governed? Are vendor, customer, and entity master data deduplicated and reconciled? Is transaction-level data accessible across the systems an agent would touch? Are data quality and ownership defined, with someone accountable? A "no" to any of these is a data-remediation task to schedule before, not during, the first deployment.

7 Governance, Risk & the EU AI Act

Governance is the deciding factor between an agent that scales and a liability that stalls. Nearly half of finance leaders will not deploy an agent without clear governance, and those who build it in from the start see faster adoption, not slower.^{[13][3]} In 2026 governance is also a legal obligation, not only good practice.

7.1 The finance AI governance stack

A workable governance model layers six controls, from inventory through monitoring. Each maps to both operational good sense and, for EU-exposed organizations, to specific obligations of the EU AI Act.


FIGURE 7.1 · THE SIX-LAYER FINANCE AI GOVERNANCE STACK

1	AI inventory and risk classification List every AI system, its use, data, and provider or deployer role; classify by risk tier.	AI Act Annex III
2	Data governance Ensure training and operating data are relevant, representative, and checked for bias and proxies.	AI Act Art. 10
3	Risk management system A documented, continuous process across the lifecycle, not a one-off assessment.	AI Act Art. 9
4	Human oversight and authorization limits Defined approval thresholds and escalation criteria; a human can understand, monitor, and override.	AI Act Art. 14
5	Logging and auditability Automatic, immutable logs detailed enough for post-hoc review; retain for the required period.	AI Act Art. 13
6	Monitoring Continuous watch for model drift, performance decay, and incidents, with prompt reporting.	Ongoing

7.2 The agent decision loop, with control built in

Governance is most durable when it is part of how the agent runs, not a review bolted on afterward. A well-designed finance agent moves through a controlled loop on every task.

FIGURE 7.2 · THE CONTROLLED AGENT DECISION LOOP



An agent that escalates an invoice above its purchase-order amount to a human approver is the loop working as designed: autonomy on the routine, human judgment on the material.^[19]

7.3 The EU AI Act: the 2026 deadline that matters

For organizations with EU exposure, the regulatory timeline is now concrete. The Act entered into force in August 2024, and its obligations for high-risk AI systems become applicable on 2 August 2026.^{[17][19]} Crucially, scope follows where the system is used, not where the firm is headquartered, so non-EU organizations with EU customers or entities can be in scope.^[18]

Several finance use cases sit squarely in the high-risk category under Annex III, notably creditworthiness assessment and credit scoring of individuals, and risk assessment and pricing in life and health insurance.^{[18][32]} For those systems, the full obligations of Articles 9 to 15 apply: a documented and continuous risk-management system, data governance with bias checks, technical documentation, automatic logging with retention, human oversight, and accuracy and cybersecurity measures.^{[18][20]} Penalties are material, reaching up to the higher of a fixed sum or a percentage of global turnover.^[20]

Risk tier	What it covers in finance	Obligation
Prohibited	A narrow set of banned practices; rarely relevant to mainstream finance.	Not permitted.
High-risk	Credit scoring, insurance pricing, evaluation of financial standing of individuals. ^[18]	Full Articles 9–15 program by Aug 2026.
Limited-risk	Chatbots and similar; many internal finance assistants.	Transparency: disclose AI interaction. ^[17]
Minimal-risk	The large majority of internal automation use cases.	No specific obligations; good practice applies. ^[20]

DO NOT GAMBLE ON THE DELAY

A Digital Omnibus package under negotiation in 2026 could conditionally move some high-risk deadlines later, but it is conditional and unconfirmed. Practitioners are uniform: treat 2 August 2026 as the binding date, because a credible readiness program for a mid-size institution takes roughly nine to fourteen months end to end, and conformity work cannot be compressed at the last minute.^{[17][19]}

7.4 Provider or deployer, and the SAP angle


Obligations differ by role. An organization that builds its own model is usually a provider and carries the heavier documentation and conformity burden; one that licenses a vendor's agent is usually a deployer and must still implement human oversight, data-governance checks, monitoring, staff training, and incident reporting.^{[18][23]} A vendor's model becomes the deployer's regulatory obligation, so contracts should require technical documentation, audit rights, change notification, and incident reporting.^[18] Embedded SAP AI capabilities such as Joule assistants and finance agents fall within this scope, which makes governance tooling and lifecycle management part of the platform decision, not an afterthought.^[21] Organizations with mature model-risk practices, for example under existing banking model-risk guidance, have a real head start on the documentation the Act expects.^[23]

8 The Business Case: ROI, Payback & KPIs

The returns are documented, but they are uneven and they reward focus. The pattern across sources is consistent: task-level automation pays back fast, enterprise-wide programs pay back over a longer horizon, and the variable that moves the timeline most is data readiness.

8.1 Payback by ambition

FIGURE 8.1 · INDICATIVE PAYBACK WINDOW BY DEPLOYMENT SCOPE (MONTHS)



Indicative ranges synthesized across vendor and advisory sources; actual payback depends heavily on data readiness and process volume.^[3]
^[4]^[11]^[14]

8.2 What the numbers say

Metric	Reported figure	Source type
Task-automation payback (for example AP)	Roughly 4 to 9 months. ^[3]	Advisory
Enterprise multi-function payback	Roughly 12 to 36 months, by scope and data readiness. ^[3]	Advisory
Autonomous-agent average ROI	Around 80%, versus about 67% for general AI projects. ^[13]	Industry survey
AP invoice-processing cost reduction	Reported 60% to 80% in some deployments. ^[11]	Vendor-reported
Agentic investment return	One study: average 2.3x within 13 months; top performers far higher. ^[4]	Advisory

Metric	Reported figure	Source type
Close-time reduction	Up to roughly 50% for early adopters. ^[8]	Vendor-reported

8.3 A KPI set that survives scrutiny

Category	Representative KPIs
Efficiency	Straight-through processing rate; cost per transaction; cycle time; headcount-to-volume ratio.
Quality & control	Exception and error rate; control-failure incidents; audit findings; percentage of decisions with complete logs.
Cash & value	Days sales outstanding; days payable optimization; close days; forecast accuracy.
Adoption & trust	Share of in-scope volume handled autonomously; escalation rate; override rate and reasons.
Governance	Coverage of AI inventory; share of high-risk systems with completed conformity work; time to detect drift.

WHERE TO APPLY SKEPTICISM

Many headline figures are vendor-reported and reflect favorable deployments. Treat them as illustrative of value categories, not as guarantees, and note the counter-evidence that many earlier AI projects returned little for lack of strategy and data.^[8] A credible business case pairs the upside with a realistic contingency and a data-readiness cost line, and measures against a pre-deployment baseline rather than a vendor benchmark.

9

Buy, Build, or Embed: The Sourcing Decision

Once the use cases and governance are set, the sourcing decision shapes cost, speed, and differentiation. The principle is simple to state and easy to get wrong: buy to accelerate standard processes, build to differentiate.^[13]

9.1 Three sourcing routes

Route	What it is	Best for	Watch-outs
Embed (in your ERP)	Agents built into the core platform, for example SAP finance agents and assistants. ^{[1][21]}	Standard processes tightly coupled to the ERP; organizations already on or moving to S/4HANA.	Tied to platform roadmap and release cadence; still your governance obligation as deployer.
Buy (specialist platform)	Best-of-breed finance agents for AP, AR, close, or treasury. ^[14]	Shared, non-differentiating processes where vendors have deep, trained models.	Integration depth and data-connector quality vary; the vendor's model is your regulatory exposure.
Build (in-house)	Custom agents on your own data and platform.	Processes that create competitive advantage, for example proprietary forecasting.	Heaviest governance and maintenance burden; you become the provider under the AI Act.

9.2 The pattern in the data

Sourcing preference splits by function in a telling way. For accounts payable, finance leaders lean toward agents embedded in existing software over building in-house; for financial planning, more lean toward self-built solutions.^[13] The logic is the buy-to-accelerate, build-to-differentiate rule in practice: a process shared across thousands of organizations is best served by a vendor that has trained on enormous volumes, while a process that is a source of competitive edge is worth owning.^[13]

A VENDOR-EVALUATION SHORT LIST

When buying or embedding, weight the criteria that actually determine success in finance: depth of integration with your ERP and data systems, explainability of agent decisions, strength of audit trails and logging, fit with your control framework, and contractual access to the documentation the AI Act requires.^{[11][18]} Feature breadth matters less than integration depth and governability.


9.3 A caution on legacy automation

Organizations carrying heavy investments in older robotic process automation should weigh maintenance honestly. Brittle, screen-based bots can consume a large share of operational capacity in upkeep and break with every interface change; reasoning-based agents are promoted partly on lower maintenance and higher straight-through processing.^[14] The migration from legacy automation to agents is itself a program to plan, not a switch to flip.

10 A 90-Day Readiness Roadmap

Readiness does not require a year before anything happens. A disciplined ninety days can produce an honest baseline, a governance framework, and a first governed agent in production. The sequence below moves from assessment to a controlled pilot, deliberately front-loading the data and governance work that determines whether the pilot scales.

FIGURE 10.1 · THE LABWYZE 90-DAY FINANCE AI READINESS ROADMAP



10.1 Principles that protect the program

- **Readiness-first, not date-first.** Commit to a go-live after the data and governance baseline, not before. The costliest surprises are data problems discovered late.^[8]
- **Pick one proving ground.** Resist the urge to automate everywhere at once. A single high-volume, structured-data use case delivers a quick win and a template.^[8]
- **Govern from day one.** Stand up the inventory, classification, and logging with the first pilot, so governance scales with the program rather than chasing it.^[19]
- **Build the team, not just the bot.** New roles, agent supervisors and value owners, and staff training are part of readiness, and training is itself an AI Act expectation for deployers.^[18]

THE ONE THING NOT TO DO

Do not run another open-ended experiment. The 2026 consensus is that boards have lost patience with pilots that do not lead anywhere; the value comes from a focused use case with real volume, measurable outcomes, controlled risk, and governance built in from the start.^{[13][11]}

11 Risks, Pitfalls & Mitigations

The failure modes for finance AI are as consistent as those for any large change program, and almost all are avoidable with the discipline already described.

Pitfall	How it shows up	Mitigation
Experiment trap	Endless pilots that test capability but never solve a business problem or scale. ^[13]	Tie every deployment to a measurable outcome with real volume; set failure criteria up front.
Poor data foundation	Agents reason on inconsistent or duplicated data and institutionalize errors at speed. ^[8]	Treat a data-readiness baseline as a prerequisite, not a parallel task.
Governance as afterthought	Agents deployed without limits, logging, or oversight; adoption stalls on risk concerns. ^[13]	Stand up the six-layer governance stack with the first pilot.
Over-autonomy too soon	High autonomy on processes the data and controls cannot yet support; fast mistakes. ^[15]	Start bounded and human-approved; earn autonomy with each audited cycle.
Regulatory blind spot	High-risk use cases run without conformity work as the August 2026 deadline passes. ^[17]	Inventory and classify early; complete Articles 9–15 work for high-risk systems on time.
Vendor risk transfer	A licensed model's gaps become the deployer's regulatory liability. ^[18]	Require documentation, audit rights, change notice, and incident reporting in contracts.
Talent and trust gap	Staff resist or misuse agents; oversight roles undefined. ^[18]	Define new roles, train supervisors, and make AI literacy part of the rollout.
Legacy automation drag	Brittle older bots consume capacity and break on every change. ^[14]	Plan the migration from legacy automation deliberately, not as a switch.

THE HONEST CAVEAT ON AGENTS

Autonomy is a double-edged capability: faster decisions, but faster mistakes if the guardrails are wrong.^[15] The most powerful protection is not a better model but better governance and a deliberately incremental rollout, the same lesson that recovered programs in every prior wave of finance technology have taught.

12 Getting Started

The organizations pulling ahead in 2026 are not the ones with the most advanced models. They are the ones that assessed honestly, fixed their data, governed from the start, and proved value on one use case before scaling. That sequence is repeatable.

12.1 The first five moves

1. **Run a readiness assessment.** Place the function on the maturity model, baseline the data foundation, and inventory existing AI by risk tier.
2. **Pick the proving ground.** Choose one high-volume, structured-data use case, usually accounts payable, with a measurable outcome.
3. **Stand up governance.** Implement the six-layer stack so it scales with the program from the first pilot.
4. **Decide sourcing deliberately.** Buy or embed for standard processes; build only where it differentiates.
5. **Address the regulatory clock.** For EU-exposed high-risk use cases, start the conformity work now against the August 2026 deadline.

HOW LABWYZE CAN HELP

labwyze works with finance leaders to turn this playbook into a plan: a structured AI maturity and data-readiness assessment for the finance function, a prioritized use-case roadmap, a governance framework mapped to the EU AI Act, and a path from first pilot to scaled, autonomous operations. The goal is a focused assessment that identifies where agentic AI delivers the fastest, safest return in your organization, and a ninety-day plan to get the first governed agent into production.

This playbook is a synthesis of public sources current to mid-2026 and is intended as planning guidance, not legal, tax, or investment advice. Regulatory requirements evolve; confirm current obligations with qualified counsel for your jurisdictions.

A

Appendix: Glossary & Sources

A.1 Glossary

Agentic AI

Autonomous systems that perceive context, plan multi-step actions, and execute end-to-end within defined rules, rather than producing output for a human to act on.

Assistant vs agent

An assistant supports a human who decides and acts; an agent decides and acts within its authority, escalating exceptions.

Autonomous Finance

A finance function in which agents run core processes continuously within guardrails, with human oversight by exception.

Continuous close

An operating model in which reconciliation and controls run continuously rather than concentrating at month-end.

Deployer / provider

Under the EU AI Act, a provider develops or places an AI system on the market; a deployer uses one under its own authority. Obligations differ.

EU AI Act

The European Union's risk-based AI regulation; high-risk obligations apply from 2 August 2026.

Human-in-the-loop

A control pattern in which a human approves material agent decisions before execution.

Straight-through processing (STP)

The share of transactions completed without manual intervention; a core efficiency KPI for finance agents.

Unified data fabric

A consolidated, transaction-level data layer across ERP, planning, and treasury that gives agents a single source of truth.

Clean core

Keeping the ERP core close to standard and pushing extensions to a platform layer, so innovation can be adopted quickly.

A.2 Sources

External sources used in this playbook, numbered as cited inline. Vendor and partner materials are labelled where relevant; independent analyst, survey, and regulatory sources are noted as such. Figures synthesized across sources are planning benchmarks.

1. SAP, "Sapphire 2026 Innovation News Guide" (2026). Autonomous Enterprise, Autonomous Finance, Joule assistants and agents, Financial Closing Assistant. sap.com
2. Deloitte, "CFO Signals" survey (Jan 2026), reported via CFO Dive. 54% AI-agent integration priority. cfodive.com
3. Lydonia, "The CFO's Guide to Agentic AI" (2026). Payback windows, sequencing, governance. lydonia.ai
4. Azilen, "Agentic AI in Financial Services: 2026 Guide" (2026). Readiness path; KPMG and McKinsey ROI figures. azilen.com
5. assistants.ai, "Agentic AI in Finance and Accounting: 2026 CFO Guide" (2026). Deloitte and Gartner projections. assistants.ai
6. ChatFin, "2026 Finance AI Spending" (2026). Copilot-to-agent shift; inverted risk calculus. chatfin.ai
7. CFO Dive, "HPE CFO puts agentic AI at center of 2026 finance priorities" (Feb 2026). Deloitte 54%; sequencing. cfodive.com

8. Auxis, "Agentic AI for Finance and Accounting" (2025). AP use cases; close-time reduction; Gartner 86% no-ROI; data foundation. auxis.com
9. Workday, "AI Agents in Finance: Use Cases" (2025). Enabling layers; close, treasury, FP&A; Gartner projections. workday.com
10. Datamatics, "Agentic AI Use Cases in Finance & Accounting" (2026). Use-case catalog by function. datamatics.com
11. Accelirate, "AI Agents in Finance: ROI and Roadmap" (2026). ROI multiples; AP cost reduction; four-phase approach. accelirate.com
12. CFA Institute Research, "Agentic AI for Finance" (2026). Workflows-versus-agents distinction; control in regulated finance. cfainstitute.org
13. Artificial Intelligence News / Basware & FT Longitude (Feb 2026). AP as proving ground; autonomous-agent ROI; buy-versus-build; governance. artificialintelligence-news.com
14. Kognitos, "Agentic AI Platforms for Finance Automation 2026" (2026). Payback and STP comparisons; ERP integration. kognitos.com
15. Houseblend, "AI Agents in Finance 2026: Reality vs Hype" (2026). Amplifier-not-replacement; incremental pilots. houseblend.io
16. Deloitte, "Tech Trends 2026: CFO Guide" (2026). Agentic workforce; high-ROI use cases. deloitte.com
17. European Commission, "AI Act: Regulatory framework" (2026). Timeline; Aug 2026 applicability; Digital Omnibus. europa.eu
18. Finextra (Borisa), "The EU AI Act's August 2026 Deadline" (May 2026). Annex III high-risk finance; Articles 9, 10, 13; scope. finextra.com
19. Alice Labs, "EU AI Act for Financial Services" (May 2026). Provider vs deployer; inventory-first; readiness duration. alicelabs.ai
20. Compliquet, "EU AI Act 2026 Requirements" (2026). Risk tiers; penalties; inventory-to-conformity path. compliquet.com
21. Arvato Systems, "EU AI Act: Compliance for the CFO Sector" (May 2026). SAP Joule and finance agents in scope; governance tooling. arvato-systems.com
22. Auritas, "2026 SAP Trends" (Mar 2026). Finance automation priorities; data readiness. auritas.com
23. FluxForce, "EU AI Act Article 6: High-Risk Requirements" (2026). Provider/deployer burden; model-risk head start; supervisory focus. fluxforce.ai
24. SNP Group, "Key Trends Shaping SAP Transformations in 2026" (Dec 2025). Data foundation for AI; fix-during-migration. snpgroup.com

A NOTE ON EVIDENCE QUALITY

Regulatory facts draw on official EU sources and independent legal analysis. Adoption, ROI, and outcome figures are benchmarks synthesized across vendor, advisory, and survey sources and vary widely; they should inform a readiness-based estimate for your organization, not substitute for one. Vendor-reported outcomes are illustrative. Where a claim depends on a single or partisan source, that limitation is flagged in the body.